


Semi-supervised Classification with Graph Convolutional Neural Networks

by Thomas N. Kipf and Max Welling

Daan Michiels

February 21, 2018

Semi-supervised classification on graphs



undirected, connected
 N nodes
each node has C features

some nodes are labeled
want to label the rest

Existing approach:
smoothing labels over graph

$$\mathcal{L} = \mathcal{L}_0 + \lambda \mathcal{L}_{\text{reg}}$$

Supervised loss

$$\sum_{i,j} A_{ij} \|f(x_i) - f(x_j)\|^2$$

Idea:
CNN on graphs
(filtering? convolution?)

Proposal: spectral convolutions, to first order (kind of)

Convolutions on graphs: expensive $O(N^2)$

→ Approximate, in this case “to first order”
(in Fourier domain, multiply by linear function)

Important property: only need a node and its neighbors!

After *renormalization trick*, get

e.g. ReLU

$$H^{(l+1)} = \sigma \left(\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H^{(l)} W^{(l)} \right)$$

activation of layer l

trainable weights

and $H^{(0)} = X$, the node features.

$$\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$$

$$\tilde{A} = A + I_N$$

Classification setup

$N \times N$

Calculate $\hat{A} = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$ in pre-processing.

Two-layer model for classification:

$$Z = \text{softmax} \left(\hat{A} \text{ReLU} \left(\hat{A} X W^{(0)} \right) W^{(1)} \right).$$

Time complexity: $O(|\mathcal{E}|CHF)$

$N \times C$


$C \times H$

$H \times F$

Batch gradient descent

Dropout

TensorFlow on GPU



Experiments

Three types: citation networks, knowledge graph, random graphs

Table 2: Summary of results in terms of classification accuracy (in percent).

Method	Citeseer	Cora	Pubmed	NELL
ManiReg [3]	60.1	59.5	70.7	21.8
SemiEmb [28]	59.6	59.0	71.1	26.7
LP [32]	45.3	68.0	63.0	26.5
DeepWalk [22]	43.2	67.2	65.3	58.1
ICA [18]	69.1	75.1	73.9	23.1
Planetoid* [29]	64.7 (26s)	75.7 (13s)	77.2 (25s)	61.9 (185s)
GCN (this paper)	70.3 (7s)	81.5 (4s)	79.0 (38s)	66.0 (48s)
GCN (rand. splits)	67.9 \pm 0.5	80.1 \pm 0.5	78.9 \pm 0.7	58.4 \pm 1.7

Experiments

Table 3: Comparison of propagation models.

Description	Propagation model	Citeseer	Cora	Pubmed
Chebyshev filter (Eq. 5)	$K = 3$	69.8	79.5	74.4
	$K = 2$	69.6	81.2	73.8
1 st -order model (Eq. 6)	$X\Theta_0 + D^{-\frac{1}{2}}AD^{-\frac{1}{2}}X\Theta_1$	68.3	80.0	77.5
Single parameter (Eq. 7)	$(I_N + D^{-\frac{1}{2}}AD^{-\frac{1}{2}})X\Theta$	69.3	79.2	77.4
Renormalization trick (Eq. 8)	$\tilde{D}^{-\frac{1}{2}}\tilde{A}\tilde{D}^{-\frac{1}{2}}X\Theta$	70.3	81.5	79.0
1 st -order term only	$D^{-\frac{1}{2}}AD^{-\frac{1}{2}}X\Theta$	68.7	80.5	77.8
Multi-layer perceptron	$X\Theta$	46.5	55.1	71.4




Figure 2: Wall-clock time per epoch for random graphs. (*) indicates out-of-memory error.

Limitations and future work

- ▶ Memory requirement
- ▶ No directed graphs, edge features
- ▶ Assumption of locality
- ▶ Importance of self-connections vs. edges? ($\tilde{A} = A + \lambda I_N$)
- ▶ Poor performance for highly regular graphs (e.g. image classification)

Reference: arXiv:1609.02907

Questions?